

eda

December 8, 2025

1 Exoplanet Habitability Analysis Pipeline

```
[1]: import sys
from pathlib import Path

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans, DBSCAN
from sklearn.metrics import silhouette_score
import shap

/opt/anaconda3/envs/data512/lib/python3.14/site-packages/tqdm/auto.py:21:
TqdmWarning: IPython not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
    from .autonotebook import tqdm as notebook_tqdm

[2]: if Path.cwd().name == "notebooks":
    project_root = Path.cwd().parent
    sys.path.insert(0, str(project_root))
else:
    project_root = Path.cwd()
    sys.path.insert(0, str(project_root))

# Data directories (relative paths)
data_raw = project_root / "data" / "raw"
data_processed = project_root / "data" / "processed"
reports_dir = project_root / "reports"

# Create directories if needed
data_raw.mkdir(parents=True, exist_ok=True)
data_processed.mkdir(parents=True, exist_ok=True)
reports_dir.mkdir(parents=True, exist_ok=True)

[3]: from src.build_features import add_derived_columns, DERIVED_COLUMNS
from src.impute_dataset import impute
```

```

from src.feature_selection import compute_feature_importance, □
    ↪compute_extended_esi as compute_extended_esi_rf, ESI_FEATURES
from src.esi import compute_esi, compute_esi_radius_mass_only, □
    ↪compute_extended_esi

```

[4]: # Set style
`sns.set_style("whitegrid")
plt.rcParams["figure.figsize"] = (12, 8)`

[5]: raw_path = data_raw / "exoplanets_pscomppars.csv"
processed_path = data_processed / "processed_exoplanets.csv"
imputed_path = data_processed / "processed_exoplanets_imputed.csv"

1.1 Phase 1

1.1.1 Data Ingestion

Note: Before running this notebook, please run `notebooks/data_ingestion.ipynb` first to download and save the raw data files.

[6]: `if not raw_path.exists():
 raise FileNotFoundError(
 f"Raw data not found at {raw_path}. "
 f"Please run notebooks/data_ingestion.ipynb first to download and save"
 ↪the data."
)

df_raw = pd.read_csv(raw_path)
print(f"Loaded raw data. Shape: {df_raw.shape}")

earth_in_raw = df_raw["pl_name"].str.lower() == "earth"
if not earth_in_raw.any():
 print("Warning: Earth not found in raw data. Data may need to be"
 ↪re-downloaded.")

df_raw.head()`

Loaded raw data. Shape: (6060, 683)

	objectid	pl_name	pl_letter	hostid	hostname	hd_name	hip_name	\
0	3.2390	Kepler-1167	b	2.136990	Kepler-1167	NaN	NaN	
1	3.1444	Kepler-1740	b	2.433343	Kepler-1740	NaN	NaN	
2	3.4135	Kepler-1581	b	2.442550	Kepler-1581	NaN	NaN	
3	3.6590	Kepler-644	b	2.512738	Kepler-644	NaN	NaN	
4	3.1575	Kepler-1752	b	2.507010	Kepler-1752	NaN	NaN	

	tic_id	disc_pubdate	disc_year	... cb_flag	pl_angsep	pl_angseperr1	\
0	TIC 273875149	2016-05	2016.0	...	0.0	0.0213	NaN

```

1 TIC 138479461      2022-02      2021.0 ...    0.0    0.0734      NaN
2 TIC 121215710      2016-05      2016.0 ...    0.0    0.1390      NaN
3 TIC 271669616      2016-05      2016.0 ...    0.0    0.0352      NaN
4 TIC 417655835      2022-02      2021.0 ...    0.0    0.2800      NaN

pl_angseperr2 pl_angseplim pl_angsepformat pl_angsepstr  pl_angsepsymerr \
0      NaN        0.0        NaN        0.0213      NaN
1      NaN        0.0        NaN        0.0734      NaN
2      NaN        0.0        NaN        0.1390      NaN
3      NaN        0.0        NaN        0.0352      NaN
4      NaN        0.0        NaN        0.2800      NaN

pl_angsep_reflink  pl_ndispec
0 <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
1 <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
2 <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
3 <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
4 <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0

[5 rows x 683 columns]

```

1.1.2 Build Derived Features

Compute research-backed derived features:

- Surface gravity & escape velocity (Schulze-Makuch et al., 2011)
- Bulk density (Seager et al., 2007)
- Habitable zone bounds (Kasting et al., 1993)
- Eccentricity habitability score (Rodríguez-Mozos & Moya, 2025)

```
[7]: if processed_path.exists():
    df_processed = pd.read_csv(processed_path)
    print(f"Loaded processed data. Shape: {df_processed.shape}")
else:
    print("Building derived features...")
    df_processed = add_derived_columns(df_raw.copy())
    df_processed.to_csv(processed_path, index=False)
    print(f"Saved processed dataset: {df_processed.shape}")

print(f"\nDerived features created: {[c for c in DERIVED_COLUMNS if c in df_processed.columns]}")
df_processed.head()
```

Loaded processed data. Shape: (6060, 691)

Derived features created: ['pl_surfgrav_m_s2', 'pl_escvel_km_s', 'pl_dens_calc', 'st_luminosity_solar', 'hz_inner_au', 'hz_outer_au', 'in_habitable_zone']

```
[7]: objectid      pl_name pl_letter      hostid      hostname hd_name hip_name \
0   3.2390  Kepler-1167 b          b  2.136990  Kepler-1167      NaN      NaN
1   3.1444  Kepler-1740 b          b  2.433343  Kepler-1740      NaN      NaN
```

```

2    3.4135  Kepler-1581 b      b  2.442550  Kepler-1581      NaN      NaN
3    3.6590  Kepler-644 b      b  2.512738  Kepler-644      NaN      NaN
4    3.1575  Kepler-1752 b      b  2.507010  Kepler-1752      NaN      NaN

          tic_id disc_pubdate disc_year ... \
0  TIC 273875149      2016-05  2016.0 ...
1  TIC 138479461      2022-02  2021.0 ...
2  TIC 121215710      2016-05  2016.0 ...
3  TIC 271669616      2016-05  2016.0 ...
4  TIC 417655835      2022-02  2021.0 ...

          pl_angsep_reflink pl_ndispec \
0  <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
1  <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
2  <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
3  <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0
4  <a refstr=CALCULATED_VALUE href=/docs/pscp_cal...      0.0

          pl_surfgrav_m_s2 pl_escvel_km_s pl_dens_calc st_luminosity_solar \
0            NaN          NaN          NaN        0.308168
1            NaN          NaN          NaN        0.778412
2            NaN          NaN          NaN       1.785102
3            NaN          NaN          NaN       6.091017
4            NaN          NaN          NaN       0.531969

          hz_inner_au hz_outer_au in_habitable_zone ecc_hab_score
0    0.529294     0.762528           0         1.0
1    0.841218     1.211900           0         1.0
2    1.273900     1.835243           0         1.0
3    2.353144     3.390057           0         1.0
4    0.695420     1.001856           0         1.0

[5 rows x 691 columns]

```

1.1.3 Impute Missing Values

Apply two-stage imputation: 1. **Deterministic reconstruction:** Mass-radius relations (Chen & Kipping, 2017) 2. **Iterative imputation:** MICE with predictive mean matching (van Buuren, 2018) - Excludes Earth/Solar System rows from model fitting - Tracks imputation flags for uncertainty propagation

Then **calculate ESI** after imputation to ensure all planets have complete data for ESI computation.

Note: Row filtering happens later during feature selection based on required core features.

```
[8]: if imputed_path.exists():
    df_imputed = pd.read_csv(imputed_path)
    print(f"Loaded imputed data. Shape: {df_imputed.shape}")
```

```

else:
    print(f"Applying imputation on df with shape {df_processed.shape} (this may take some time)...")
    df_imputed = impute(df_processed)
    df_imputed.to_csv(imputed_path, index=False)
    print(f"Imputed dataset saved: {df_imputed.shape}")

imputation_flags = [c for c in df_imputed.columns if "_flag" in c or "_imputed_flag" in c]
print(f"\nImputation flags: {imputation_flags[:5]}... ({len(imputation_flags)} total)")

if imputation_flags:
    for flag in imputation_flags[:5]:
        imputed_count = int(df_imputed[flag].sum())
        print(f"\t{flag}: {imputed_count}/{len(df_imputed)} values imputed")

```

Loaded imputed data. Shape: (6060, 704)

Imputation flags: ['pl_controv_flag', 'dkin_flag', 'ttv_flag', 'ptv_flag', 'tran_flag']... (26 total)

- pl_controv_flag: 40/6060 values imputed
- dkin_flag: 1/6060 values imputed
- ttv_flag: 478/6060 values imputed
- ptv_flag: 2/6060 values imputed
- tran_flag: 4504/6060 values imputed

```
[9]: # Calculate ESI after imputation (on complete data)
print("\nCalculating ESI on imputed data...")
df_imputed["esi"] = compute_esi(df_imputed)
df_imputed["esi_radius_mass"] = compute_esi_radius_mass_only(df_imputed)

print(f"ESI statistics: mean={df_imputed['esi'].mean():.4f}, median={df_imputed['esi'].median():.4f}")
df_imputed.head()
```

Calculating ESI on imputed data...

ESI statistics: mean=0.3032, median=0.3082

	objectid	pl_name	pl_letter	hostid	hostname	hd_name	hip_name	
0	3.2390	Kepler-1167	b	2.136990	Kepler-1167	NaN	NaN	
1	3.1444	Kepler-1740	b	2.433343	Kepler-1740	NaN	NaN	
2	3.4135	Kepler-1581	b	2.442550	Kepler-1581	NaN	NaN	
3	3.6590	Kepler-644	b	2.512738	Kepler-644	NaN	NaN	
4	3.1575	Kepler-1752	b	2.507010	Kepler-1752	NaN	NaN	

	tic_id	disc_pubdate	disc_year	pl_orbper_mice_flag	
0	TIC 273875149	2016-05	2016.0	...	0

```

1 TIC 138479461      2022-02    2021.0 ...          0
2 TIC 121215710      2016-05    2016.0 ...          0
3 TIC 271669616      2016-05    2016.0 ...          0
4 TIC 417655835      2022-02    2021.0 ...          0

pl_dens_mice_flag st_teff_mice_flag st_met_mice_flag st_mass_mice_flag \
0             0             0             0             0
1             0             0             0             0
2             0             0             0             0
3             0             0             0             0
4             0             0             0             0

st_rad_mice_flag pl_surfgrav_m_s2_mice_flag pl_escvel_km_s_mice_flag \
0             0             1             1
1             0             1             1
2             0             1             1
3             0             1             1
4             0             1             1

esi  esi_radius_mass
0  0.000000      0.394898
1  0.294312      0.226551
2  0.000000      0.587326
3  0.000000      0.200661
4  0.280907      0.202688

[5 rows x 706 columns]

```

1.2 Phase 2: What ESI Misses

1.2.1 Feature Importance via Random Forest + SHAP (Excluding ESI Features)

Key Innovation: ESI is defined by 5 physical parameters (radius, mass, density, escape velocity, temperature). However, habitability depends on many other factors!

We train a Random Forest to predict ESI using **ONLY non-ESI features** (orbital characteristics, stellar properties, derived metrics). This reveals: 1. What factors beyond ESI's definition predict Earth-like conditions 2. Whether the ESI metric is incomplete 3. Which orbital/stellar features matter for habitability

If R^2 is high (>0.7), it proves ESI is incomplete - other features strongly predict habitability

```
[ ]: print("Computing feature importance using ESI as target...")
```

```

results = compute_feature_importance(
    df_imputed.copy(),
    output_dir=reports_dir,
    correlation_threshold=0.90,
    min_features_required=0.5,
```

```

)

importance_df = results.importance_df
feature_matrix = results.feature_matrix
X = results.X
model = results.model
shap_values = results.shap_values
df_filtered = results.df_filtered

# Use filtered dataframe for all subsequent analysis
df_imputed = df_filtered.copy()
print(f"\nFiltered dataset: {len(df_imputed)} planets")

print(f"Feature matrix shape: {feature_matrix.shape}")
print(f"Standardized features shape: {X.shape}")

earth_mask = df_imputed["pl_name"].str.lower() == "earth"
if not earth_mask.any():
    raise ValueError("Earth not found in dataset!")

```

Computing feature importance using ESI as target...
 Filtered to 6051 planets with required core features: ['pl_rade', 'pl_masse']

Found 3 highly correlated feature pairs (>=0.9):
 pl_orbeccen <-> ecc_hab_score: 0.986
 pl_insol <-> pl_orbper: 0.962
 pl_orbsmax <-> pl_orbper: 0.903

Removing 2 highly correlated features: ['pl_orbeccen', 'pl_orbper']

[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
 [Parallel(n_jobs=-1)]: Done 34 tasks | elapsed: 0.4s
 [Parallel(n_jobs=-1)]: Done 184 tasks | elapsed: 1.6s
 [Parallel(n_jobs=-1)]: Done 434 tasks | elapsed: 3.3s
 [Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed: 3.7s finished
 [Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
 [Parallel(n_jobs=8)]: Done 34 tasks | elapsed: 0.0s
 [Parallel(n_jobs=8)]: Done 184 tasks | elapsed: 0.0s
 [Parallel(n_jobs=8)]: Done 434 tasks | elapsed: 0.1s
 [Parallel(n_jobs=8)]: Done 500 out of 500 | elapsed: 0.1s finished

R² (non-ESI → ESI): 0.5105

Top 10 features predicting ESI (excluding ESI-defining features):

feature	mean_abs_shap
pl_orbsmax	0.044483
st_rad	0.022091
ecc_hab_score	0.021895

```

pl_surfgrav_m_s2      0.020958
    st_met          0.014478
    st_mass          0.013208
    pl_insol         0.011916
    st_teff          0.004490
in_habitable_zone     0.000048

```

Filtered dataset: 6,051 planets
 Feature matrix shape: (6051, 15)
 Standardized features shape: (6051, 14)

Top 10 features by SHAP importance (for predicting ESI):

	feature	mean_abs_shap
0	pl_orbsmax	0.044483
1	st_rad	0.022091
2	ecc_hab_score	0.021895
3	pl_surfgrav_m_s2	0.020958
4	st_met	0.014478
5	st_mass	0.013208
6	pl_insol	0.011916
7	st_teff	0.004490
8	in_habitable_zone	0.000048

1.2.2 Extended ESI: Combining Traditional + Data-Driven Features

Goal: Create “Proposed ESI” by training RF on ESI features + top N non-ESI features. This tests if adding new features improves predictive power.

```
[11]: N_EXTENDED_FEATURES = 10
top_non_esi = importance_df.head(N_EXTENDED_FEATURES) ['feature'].tolist()

print(f"Training Extended ESI RF model ({len(ESI_FEATURES)} ESI +"
      f"{N_EXTENDED_FEATURES} non-ESI features)...")
extended_results = compute_extended_esi_rf(feature_matrix, top_non_esi,
                                             N_EXTENDED_FEATURES)

print("\nFeature importance in Extended model:")
print(extended_results.importance_df.to_string(index=False))
```

Training Extended ESI RF model (5 ESI + 10 non-ESI features)...

```
[Parallel(n_jobs=-1)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done  34 tasks      | elapsed:   0.4s
[Parallel(n_jobs=-1)]: Done 184 tasks      | elapsed:   2.1s
[Parallel(n_jobs=-1)]: Done 434 tasks      | elapsed:   4.9s
[Parallel(n_jobs=-1)]: Done 500 out of 500 | elapsed:   6.0s finished
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
[Parallel(n_jobs=8)]: Done  34 tasks      | elapsed:   0.0s
[Parallel(n_jobs=8)]: Done 184 tasks      | elapsed:   0.1s
```

```
[Parallel(n_jobs=8)]: Done 434 tasks      | elapsed:  0.1s
[Parallel(n_jobs=8)]: Done 500 out of 500 | elapsed:  0.1s finished
```

Extended ESI model (14 features):

ESI features: 5

Non-ESI features: 10

R² (extended → ESI): 0.9856

Feature importance in Extended model:

feature	mean_abs_shap
pl_masse	1.169419e-01
pl_escvel_km_s	7.340136e-02
pl_rade	1.789473e-02
pl_eqt	1.126638e-02
pl_dens_calc	3.491513e-03
st_teff	1.100165e-03
st_mass	1.067422e-03
pl_insol	9.842265e-04
pl_orbsmax	3.335343e-04
st_rad	3.052426e-04
pl_surfgrav_m_s2	2.806430e-04
ecc_hab_score	1.021848e-04
st_met	8.271625e-05
in_habitable_zone	3.463172e-07

```
[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.
```

```
[Parallel(n_jobs=8)]: Done 34 tasks      | elapsed:  0.0s
```

```
[Parallel(n_jobs=8)]: Done 184 tasks      | elapsed:  0.0s
```

```
[Parallel(n_jobs=8)]: Done 434 tasks      | elapsed:  0.1s
```

```
[Parallel(n_jobs=8)]: Done 500 out of 500 | elapsed:  0.1s finished
```

```
[12]: # Compute Extended ESI using formula (geometric mean like original ESI)
earth_idx = df_imputed[earth_mask].index[0]
df_imputed['extended_esi'] = compute_extended_esi(df_imputed, top_non_esi, ↴
                                                earth_idx)

# Also compute RF-based prediction for comparison
X_extended = feature_matrix[extended_results.extended_features].copy()
df_imputed['rf_prediction'] = extended_results.model.predict(X_extended)

# Create rankings
df_imputed['rank_esi'] = df_imputed['esi'].rank(ascending=False, method='min')
df_imputed['rank_extended'] = df_imputed['extended_esi'].rank(ascending=False, ↴
                                                               method='min')
df_imputed['rank_shift'] = df_imputed['rank_esi'] - df_imputed['rank_extended']

print("Extended ESI calculated using geometric mean formula")
```

```

print(f"\tOriginal ESI (5 features): [{df_imputed['esi'].min():.4f},\n
    ↪{df_imputed['esi'].max():.4f}]")
print(f"\tExtended ESI ({5+N_EXTENDED_FEATURES} features):\n
    ↪[{df_imputed['extended_esi'].min():.4f}, {df_imputed['extended_esi'].max():.
    ↪4f}]")
print(f"\nCorrelation: Extended ESI vs Original ESI = {df_imputed['esi'].
    ↪corr(df_imputed['extended_esi']):.3f}")
print(f"Correlation: Extended ESI vs RF prediction =\n
    ↪{df_imputed['extended_esi'].corr(df_imputed['rf_prediction']):.3f}")

```

[Parallel(n_jobs=8)]: Using backend ThreadingBackend with 8 concurrent workers.

[Parallel(n_jobs=8)]: Done 34 tasks | elapsed: 0.0s

[Parallel(n_jobs=8)]: Done 184 tasks | elapsed: 0.0s

Extended ESI calculated using geometric mean formula

Original ESI (5 features): [0.0000, 0.9999]

Extended ESI (15 features): [0.0000, 1.0000]

Correlation: Extended ESI vs Original ESI = 0.141

Correlation: Extended ESI vs RF prediction = 0.142

[Parallel(n_jobs=8)]: Done 434 tasks | elapsed: 0.1s

[Parallel(n_jobs=8)]: Done 500 out of 500 | elapsed: 0.1s finished

1.2.3 Ranking Comparison: Original ESI vs Proposed ESI

```

[13]: print("=*100)
print("RANKING COMPARISON: Original ESI vs Extended ESI")
print("=*100)

# Overlap analysis
top_20_original = set(df_imputed.nlargest(20, 'esi')['pl_name'])
top_20_proposed = set(df_imputed.nlargest(20, 'extended_esi')['pl_name'])
overlap = top_20_original & top_20_proposed
new_candidates = top_20_proposed - top_20_original
dropped = top_20_original - top_20_proposed

print(f"\nTop 20 Overlap: {len(overlap)}/20 ({len(overlap)/20*100:.0f}%)")
print(f"New candidates (Extended ESI): {len(new_candidates)}")
print(f"Dropped (from original ESI): {len(dropped)}")

# Biggest movers
cols_display = ['pl_name', 'rank_esi', 'rank_extended', 'rank_shift', 'esi',\n    ↪'extended_esi'] + top_non_esi[:3]

print("\nBiggest GAINERS (improved rank with Extended ESI):")
gainers = df_imputed.nlargest(10, 'rank_shift')[cols_display]
print(gainers.to_string(index=False))

```

```

print("\nBiggest LOSERS (dropped rank with Extended ESI):")
losers = df_imputed.nsmallest(10, 'rank_shift')[cols_display]
print(losers.to_string(index=False))

print("\nNew candidates in Top 20 Extended ESI:")
if new_candidates:
    for planet in list(new_candidates)[:10]:
        row = df_imputed[df_imputed['pl_name'] == planet].iloc[0]
        print(f" {planet:25s} Rank: {int(row['rank_extended']):4d} (was"
             f"{int(row['rank_esi']):4d})")

```

RANKING COMPARISON: Original ESI vs Extended ESI

Top 20 Overlap: 2/20 (10%)

New candidates (Extended ESI): 18

Dropped (from original ESI): 18

Biggest GAINERS (improved rank with Extended ESI):

	pl_name	rank_esi	rank_extended	rank_shift	esi	extended_esi
pl_orbsmax	st_rad	ecc_hab_score				
TOI-1736	c	4816.0	709.0	4107.0	0.101468	0.317258
1.3700	1.42860		0.730727			
Kepler-553	c	4621.0	660.0	3961.0	0.111385	0.323169
0.8980	0.90200		0.742942			
TOI-2180	b	4259.0	362.0	3897.0	0.136942	0.375654
0.8280	1.63600		0.730834			
Kepler-1514	b	4525.0	688.0	3837.0	0.115821	0.319676
0.7530	1.28900		0.713776			
TOI-2010	b	4166.0	346.0	3820.0	0.146939	0.377446
0.5516	1.07900		0.825083			
Kepler-1704	b	4347.0	549.0	3798.0	0.128712	0.341178
2.0270	1.69700		0.520833			
TIC 393818343	b	4655.0	862.0	3793.0	0.109469	0.295280
0.1291	1.08600		0.622743			
HD 191806	b	5127.0	1359.0	3768.0	0.080006	0.221172
2.7490	1.40223		0.824402			
TOI-4562	b	4278.0	516.0	3762.0	0.135140	0.346236
0.7680	1.15200		0.568182			
HD 16175	b	4914.0	1160.0	3754.0	0.096457	0.249913
2.1200	1.72000		0.625000			

Biggest LOSERS (dropped rank with Extended ESI):

pl_name	rank_esi	rank_extended	rank_shift	esi
---------	----------	---------------	------------	-----

extended_esi	pl_orbsmax	st_rad	ecc_hab_score			
	AU Mic d	3.0	1999.0	-1996.0	0.946443	
0.0	219.114146	0.744000	0.996959			
	Teegarden's Star b	4.0	1999.0	-1995.0	0.934802	
0.0	0.025900	0.120000	0.970874			
	OGLE-2016-BLG-0007L b	7.0	1999.0	-1992.0	0.923465	
0.0	10.100000	3.132700	NaN			
	LP 791-18 d	9.0	1999.0	-1990.0	0.913450	
0.0	0.019920	0.182000	0.998502			
	K2-384 b	11.0	1999.0	-1988.0	0.892279	
0.0	-326.559296	0.348000	NaN			
	GJ 1002 b	12.0	1999.0	-1987.0	0.891062	
0.0	0.045700	0.137000	NaN			
	Kepler-220 d	13.0	1999.0	-1986.0	0.888911	
0.0	0.163000	0.666000	1.000000			
	Kepler-1178 b	14.0	1999.0	-1985.0	0.876029	
0.0	0.159600	0.750000	1.000000			
	K2-209 b	15.0	1999.0	-1984.0	0.867328	
0.0	-259.268705	0.712365	NaN			
	LHS 475 b	16.0	1999.0	-1983.0	0.866884	
0.0	0.020370	0.278900	1.000000			

New candidates in Top 20 Extended ESI:

KOI-1783.02	Rank: 20 (was 2816)
Kepler-452 b	Rank: 6 (was 738)
Kepler-1853 b	Rank: 18 (was 506)
Kepler-139 c	Rank: 5 (was 1649)
Kepler-49 e	Rank: 19 (was 96)
Kepler-345 c	Rank: 15 (was 84)
KOI-1831 d	Rank: 4 (was 76)
K2-138 g	Rank: 17 (was 1097)
Kepler-48 d	Rank: 7 (was 516)
Kepler-725 c	Rank: 3 (was 1145)

1.2.4 Clustering on Extended Features

Find optimal k and perform K-means/DBSCAN clustering on the 15 extended features.

```
[14]: # Standardize extended features
X_extended_std = StandardScaler().fit_transform(X_extended)

# Find optimal k using silhouette score
silhouette_scores = []
for k in range(2, 11):
    kmeans_temp = KMeans(n_clusters=k, random_state=42, n_init=10)
    labels = kmeans_temp.fit_predict(X_extended_std)
    silhouette_scores.append(silhouette_score(X_extended_std, labels))
```

```

optimal_k = 2 + np.argmax(silhouette_scores)
print(f"Optimal k = {optimal_k} (silhouette score: {max(silhouette_scores):.3f})")

# K-means clustering
kmeans_ext = KMeans(n_clusters=optimal_k, random_state=42, n_init=10)
df_imputed['kmeans_ext'] = kmeans_ext.fit_predict(X_extended_std)
earth_cluster_ext = df_imputed.loc[earth_mask, 'kmeans_ext'].values[0]

# DBSCAN clustering
dbscan_ext = DBSCAN(eps=2.0, min_samples=5)
df_imputed['dbscan_ext'] = dbscan_ext.fit_predict(X_extended_std)
earth_dbscan_ext = df_imputed.loc[earth_mask, 'dbscan_ext'].values[0]

# Distance to Earth
earth_features_ext = X_extended_std[earth_mask.values][0]
df_imputed['dist_earth_ext'] = np.linalg.norm(X_extended_std - earth_features_ext, axis=1)

print(f"K-means: Earth in cluster {earth_cluster_ext}")
print(f"DBSCAN: Earth in cluster {earth_dbscan_ext}")

```

Optimal k = 2 (silhouette score: 0.922)

K-means: Earth in cluster 0

DBSCAN: Earth in cluster 0

```
[15]: # Overlap analysis: Earth's cluster vs Top 20 Extended ESI
earth_cluster_planets = set(
    df_imputed[df_imputed['kmeans_ext'] == earth_cluster_ext]
    .nsmallest(20, 'dist_earth_ext')['pl_name']
)

overlap_cluster_proposed = earth_cluster_planets & top_20_proposed

print("=="*100)
print("OVERLAP ANALYSIS: Clustering vs Extended ESI")
print("=="*100)
print("\nEarth's cluster (extended features, top 20 by proximity) vs Top 20 Extended ESI:")
print(f"\tOverlap: {len(overlap_cluster_proposed)}/20 ({len(overlap_cluster_proposed)}/20*100:.0f}%")
print("\nPlanets in both:")
for planet in list(overlap_cluster_proposed)[:10]:
    row = df_imputed[df_imputed['pl_name'] == planet].iloc[0]
    print(f"  {planet:25s} Rank(Proposed): {int(row['rank_extended']):4d}, Dist: {row['dist_earth_ext']:.3f}")
print("=="*100)
```

```
=====
=====
```

OVERLAP ANALYSIS: Clustering vs Extended ESI

```
=====
=====
```

Earth's cluster (extended features, top 20 by proximity) vs Top 20 Extended ESI:
Overlap: 5/20 (25%)

Planets in both:

Earth	Rank(Proposed): 1, Dist: 0.000
KOI-1831 d	Rank(Proposed): 4, Dist: 0.899
Kepler-48 d	Rank(Proposed): 7, Dist: 0.901
TOI-2088 b	Rank(Proposed): 12, Dist: 0.937
Kepler-324 c	Rank(Proposed): 11, Dist: 0.856

```
=====
=====
```

1.2.5 Visualizations: Extended ESI Analysis

```
[ ]: # Viz 1: Ranking shift scatter plot
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Left: Rank comparison
axes[0].scatter(df_imputed['rank_esi'], df_imputed['rank_extended'], alpha=0.3, s=20, color='gray')
axes[0].plot([1, df_imputed['rank_esi'].max()], [1, df_imputed['rank_esi'].max()],
            'r--', linewidth=2, label='No change')
axes[0].scatter(df_imputed.loc[earth_mask, 'rank_esi'], df_imputed.loc[earth_mask, 'rank_extended'],
                s=400, color='red', marker='*', edgecolors='black', linewidths=2, label='Earth', zorder=10)
axes[0].set_xlabel('Original ESI Rank', fontsize=12)
axes[0].set_ylabel('Extended ESI Rank', fontsize=12)
axes[0].set_title('Ranking Shifts: Original vs Extended ESI', fontsize=14, fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)
axes[0].text(0.02, 0.98, 'Above line = Improved rank', transform=axes[0].transAxes,
            verticalalignment='top', fontsize=10, bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

# Right: ESI comparison scatter
axes[1].scatter(df_imputed['esi'], df_imputed['extended_esi'], alpha=0.3, s=20, color='steelblue')
```

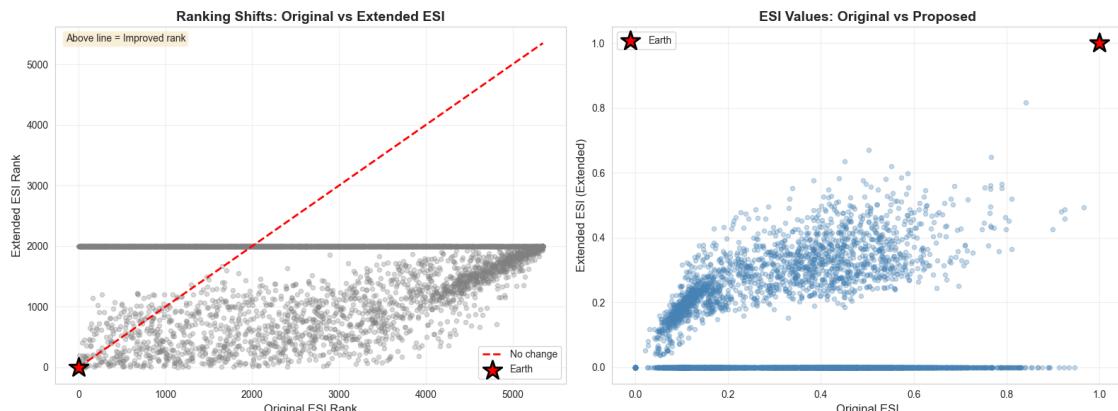
```

axes[1].scatter(df_imputed.loc[earth_mask, 'esi'], df_imputed.loc[earth_mask, 'extended_esi'],
                s=400, color='red', marker='*', edgecolors='black', linewidths=2, label='Earth', zorder=10)
axes[1].set_xlabel('Original ESI', fontsize=12)
axes[1].set_ylabel('Extended ESI (Extended)', fontsize=12)
axes[1].set_title('ESI Values: Original vs Proposed', fontsize=14, fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(reports_dir / "extended_esi_comparison.png", dpi=150, bbox_inches="tight")
plt.show()

print("Saved: extended_esi_comparison.png")

```



Saved: extended_esi_comparison.png

```

[ ]: # Viz 2: Feature contribution heatmap for top planets
top_15_proposed = df_imputed nlargest(15, 'extended_esi')
feature_values = top_15_proposed[extended_results.extended_features].T

# Normalize for better visualization
feature_values_norm = (feature_values - feature_values.min(axis=1).values[:, None]) / \
    (feature_values.max(axis=1).values[:, None] - feature_values.min(axis=1).values[:, None] + 1e-10)

fig, ax = plt.subplots(figsize=(14, 10))
im = ax.imshow(feature_values_norm, aspect='auto', cmap='RdYlGn')

```

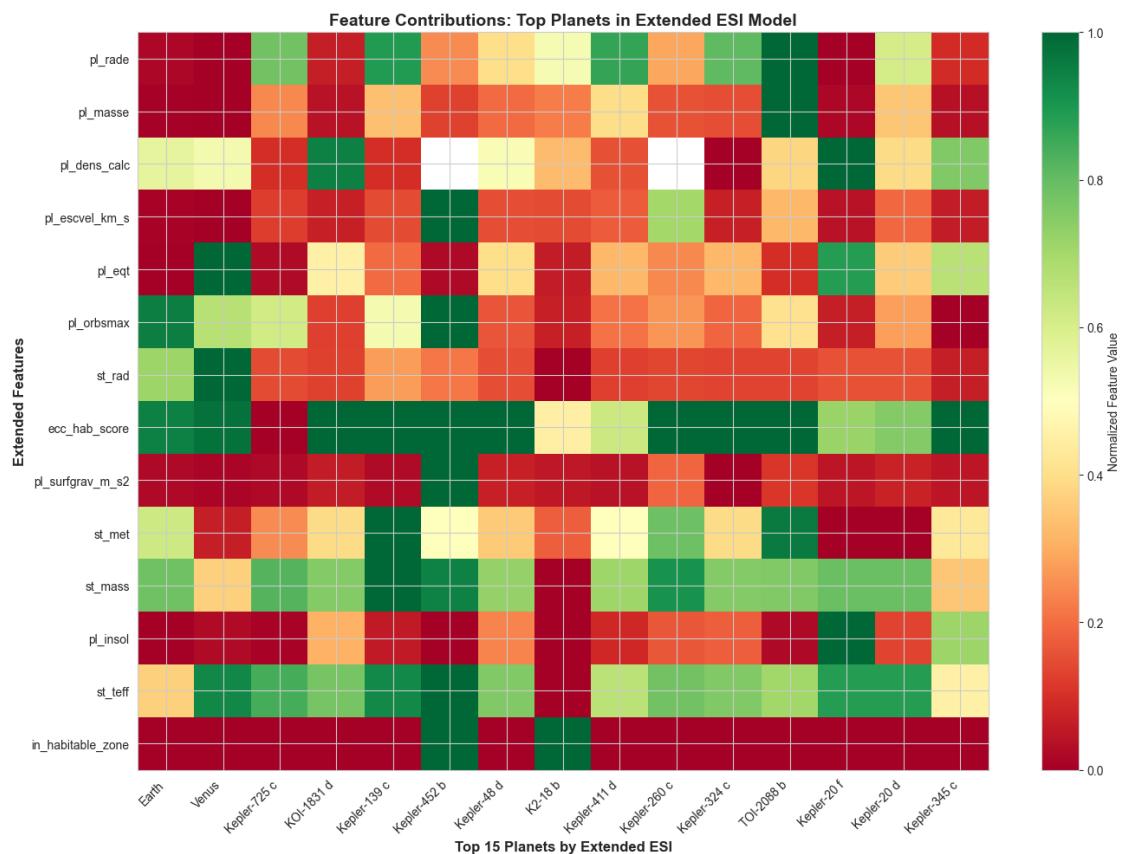
```

ax.set_xticks(range(len(top_15_proposed)))
ax.set_xticklabels(top_15_proposed['pl_name'].values, rotation=45, ha='right', fontweight='bold', fontsize=10)
ax.set_yticks(range(len(extended_results.extended_features)))
ax.set_yticklabels(extended_results.extended_features, fontsize=10)
ax.set_xlabel('Top 15 Planets by Extended ESI', fontsize=12, fontweight='bold')
ax.set_ylabel('Extended Features', fontsize=12, fontweight='bold')
ax.set_title('Feature Contributions: Top Planets in Extended ESI Model', fontsize=14, fontweight='bold')

plt.colorbar(im, ax=ax, label='Normalized Feature Value')
plt.tight_layout()
plt.savefig(reports_dir / "feature_contributions_heatmap.png", dpi=150, bbox_inches="tight")
plt.show()

print("Saved: feature_contributions_heatmap.png")

```



Saved: feature_contributions_heatmap.png

```
[18]: # Viz 3: R2 comparison bar chart
r2_values = [0.5105, extended_results.r2_score]
methods = ['Non-ESI features\nonly', f'Extended ESI\n{n_EXTENDED_FEATURES} + 5\n↳features)']

fig, ax = plt.subplots(figsize=(10, 6))
bars = ax.bar(methods, r2_values, color=['steelblue', 'darkgreen'], ↳
    edgecolor='black', linewidth=2, alpha=0.8)

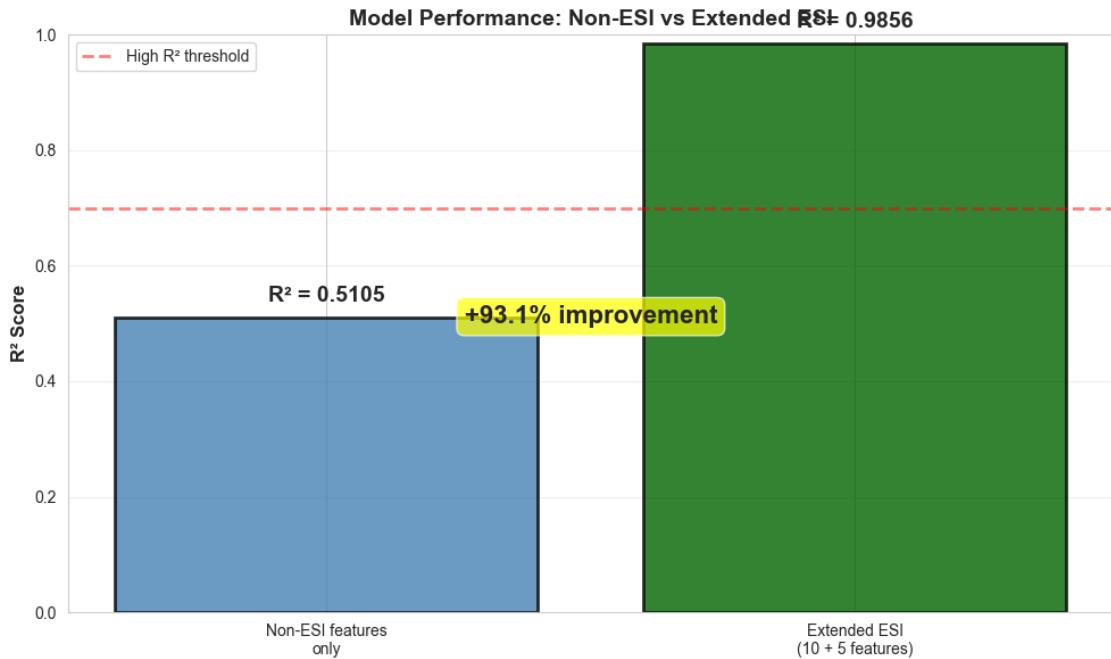
for i, (bar, val) in enumerate(zip(bars, r2_values)):
    ax.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 0.02,
            f'R2 = {val:.4f}', ha='center', va='bottom', fontsize=14, ↳
            fontweight='bold')

ax.set_ylabel('R2 Score', fontsize=12, fontweight='bold')
ax.set_title('Model Performance: Non-ESI vs Extended ESI', fontsize=14, ↳
    fontweight='bold')
ax.set_ylim(0, 1.0)
ax.grid(True, alpha=0.3, axis='y')
ax.axhline(y=0.7, color='red', linestyle='--', linewidth=2, alpha=0.5, ↳
    label='High R2 threshold')
ax.legend()

improvement = ((extended_results.r2_score - 0.5105) / 0.5105) * 100
ax.text(0.5, 0.5, f'+{improvement:.1f}% improvement', ha='center', fontsize=16,
        fontweight='bold', bbox=dict(boxstyle='round', facecolor='yellow', ↳
        alpha=0.7),
        transform=ax.transAxes)

plt.tight_layout()
plt.savefig(reports_dir / "r2_comparison.png", dpi=150, bbox_inches="tight")
plt.show()

print("Saved: r2_comparison.png")
```



Saved: r2_comparison.png

```
[19]: # Summary: Extended ESI key findings
print("\n" + "="*100)
print(" " * 30 + "EXTENDED ESI: KEY FINDINGS")
print("=".*100)

print("\n1. R^2 IMPROVEMENT:")
print("\tNon-ESI features only: R^2 = 0.5105 (51%)")
print(f"Extended ESI (15 feat): R^2 = {extended_results.r2_score:.4f} ↴({extended_results.r2_score*100:.1f}%)")
print(f"Improvement: +{((extended_results.r2_score - 0.5105)/0.5105)*100:.1f}%)"

print("\n2. TOP 3 NON-ESI FEATURES IN EXTENDED MODEL:")
for idx, row in extended_results.importance_df.head(3).iterrows():
    if row['feature'] not in ESI_FEATURES:
        print(f"\t{idx+1}. {row['feature'][:20s]} SHAP: {row['mean_abs_shap']:.4f}")

print("\n3. RANKING CHANGES:")
print(f"\tTop 20 overlap: {len(overlap)}/20 ({len(overlap)/20*100:.0f}%)")
print(f"\tNew candidates: {len(new_candidates)}")
print(f"\tDropped: {len(dropped)}")

print("\n4. CLUSTERING AGREEMENT:")
```

```

print(f"\tEarth's cluster vs Top 20 Proposed: {len(overlap_cluster_proposed)}/
    ↪20 ({len(overlap_cluster_proposed)/20*100:.0f}%)")

print("\n" + "="*100)

```

=====

=====

EXTENDED ESI: KEY FINDINGS

=====

=====

1. R² IMPROVEMENT:

Non-ESI features only: R² = 0.5105 (51%)

Extended ESI (15 feat): R² = 0.9856 (98.6%)

Improvement: +93.1%

2. TOP 3 NON-ESI FEATURES IN EXTENDED MODEL:

3. RANKING CHANGES:

Top 20 overlap:	2/20 (10%)
New candidates:	18
Dropped:	18

4. CLUSTERING AGREEMENT:

Earth's cluster vs Top 20 Proposed: 5/20 (25%)

```

=====
=====
[21]: # Use the same non-ESI feature columns that were used to compute SHAP
feature_cols = [
    col
    for col in feature_matrix.columns
    if col not in ESI_FEATURES + ["esi"]
]

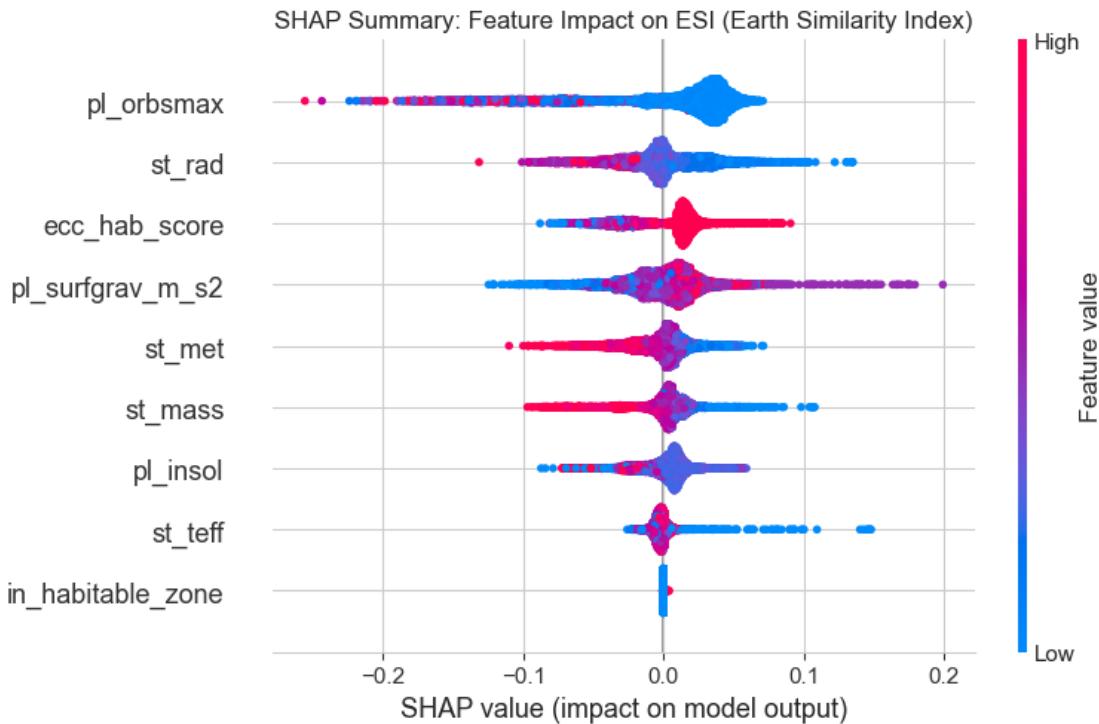
print("Generating SHAP summary plot...")
shap.summary_plot(shap_values, feature_matrix[feature_cols], show=False)
plt.title("SHAP Summary: Feature Impact on ESI (Earth Similarity Index)")
plt.tight_layout()
plt.savefig(reports_dir / "shap_detailed.png", dpi=150, bbox_inches="tight")
plt.show()

print("\nTop 20 planets by ESI (most Earth-like):")
top_esi = df_imputed.nlargest(20, "esi")[
    ["pl_name", "esi", "pl_masse", "pl_rade", "pl_eqt", "in_habitable_zone"]
]

```

```
]
print(top_esi.to_string(index=False))
```

Generating SHAP summary plot...



Top 20 planets by ESI (most Earth-like):

	pl_name	esi	pl_masse	pl_rade	pl_eqt	in_habitable_zone
	Earth	0.999936	1.000000	1.000000	255.000000	0
	TRAPPIST-1 f	0.966019	1.039000	1.045000	217.700000	0
	AU Mic d	0.946443	1.053000	1.020000	168.760873	0
	Teegarden's Star b	0.934802	1.197844	1.050000	277.000000	0
	TRAPPIST-1 e	0.925299	0.692000	0.920000	249.700000	1
	TRAPPIST-1 c	0.924315	1.308000	1.097000	339.700000	0
OGLE-2016-BLG-0007L	b	0.923465	1.320000	1.090000	193.750425	0
	TRAPPIST-1 g	0.917362	1.321000	1.129000	197.300000	0
	LP 791-18 d	0.913450	0.900000	1.032000	395.500000	0
	TRAPPIST-1 b	0.898853	1.374000	1.116000	397.600000	0
	K2-384 b	0.892279	1.311310	1.076000	326.224508	0
	GJ 1002 b	0.891062	1.115572	1.030000	230.900000	1
	Kepler-220 d	0.888911	0.927975	0.980000	401.000000	0
	Kepler-1178 b	0.876029	1.284458	1.070000	378.000000	0
	K2-209 b	0.867328	0.593544	0.868503	373.481828	0
	LHS 475 b	0.866884	0.967103	0.991000	586.000000	0

KMT-2020-BLG-0414L b	0.854582	0.960000	0.997000	812.200658	0
LT T 1445 A c	0.852195	1.540000	1.147000	508.000000	0
Venus	0.841466	0.815000	0.949900	737.000000	0
Kepler-1650 b	0.840048	0.859812	0.960000	602.000000	0

1.3 Research Question Answers

Using Extended ESI methodology to answer research questions from project proposal.

```
[22]: # RQ1: Earth's percentile position
key_features = ["pl_masse", "pl_rade", "pl_eqt", "pl_insol", "pl_escvel_km_s", ↴"pl_dens_calc"]

earth_percentiles = {}
for feat in key_features:
    if feat in df_imputed.columns:
        data = df_imputed[feat].dropna()
        earth_val = df_imputed.loc[earth_mask, feat].values[0]
        if not np.isnan(earth_val):
            pct = (data < earth_val).sum() / len(data) * 100
            earth_percentiles[feat] = (earth_val, pct)

print("*"*100)
print(" "*30 + "RESEARCH QUESTION 1")
print("How does Earth compare to other planets across key characteristics?")
print("*"*100)
print("\nEarth's Percentile Position:")
for feat, (val, pct) in earth_percentiles.items():
    print(f" {feat:20s}: {pct:5.1f}th percentile")
print("*"*100)
```

```
=====
=====
                    RESEARCH QUESTION 1
How does Earth compare to other planets across key characteristics?
=====
```

Earth's Percentile Position:

pl_masse	:	3.5th percentile
pl_rade	:	3.8th percentile
pl_eqt	:	4.3th percentile
pl_insol	:	12.3th percentile
pl_escvel_km_s	:	15.0th percentile
pl_dens_calc	:	75.0th percentile

```
[27]: # RQ2: Most similar planets using Extended ESI
earth_cluster_ext_df = df_imputed[df_imputed['kmeans_ext'] ==_
    ↪earth_cluster_ext].copy()
earth_cluster_ext_df = earth_cluster_ext_df.sort_values('extended_esi',_
    ↪ascending=False)

print("\n" + "="*100)
print(" "*30 + "RESEARCH QUESTION 2")
print("Which planets are most similar to Earth (Extended ESI approach)?")
print("=".*100)

print("\nMETHODOLOGY:")
print("\t1. Trained RF on non-ESI features → R2 = 51% (ESI incomplete)")
print(f"\t2. Trained RF on ESI + {N_EXTENDED_FEATURES} non-ESI features → R2 =_
    ↪{extended_results.r2_score:.1%}")
print("\t3. Used Extended model predictions as Extended ESI")
print("\t4. K-means clustering on 15 extended features")
print("\t5. Ranked by distance to Earth in extended feature space")

print("\nTOP 10 MOST SIMILAR PLANETS (Earth's cluster, Extended ESI):")
print(" "*120)
print(f"[{'Rank':<6}{'Planet':<25}{'Dist':<10}{'Extended ESI':<13}{'Original_'
    ↪'ESI':<13}{'Orig Rank':<11}{'Ext Rank':<11}{'Rank Shift':<11}]")
print(" "*120)

for idx, row in enumerate(earth_cluster_ext_df.head(10).itertuples(), 1):
    shift = int(row.rank_esi - row.rank_extended)
    shift_str = f"+{shift}" if shift > 0 else str(shift)
    print(
        f"[{idx:<6}{row.pl_name:<25}{row.dist_earth_ext:<10.3f}"
        f"{row.extended_esi:<13.4f}{row.esi:<13.4f}"
        f"{int(row.rank_esi):<11}{int(row.rank_extended):<11}{shift_str:<11}]"
    )

print("=".*120)
```

```
=====
=====
                    RESEARCH QUESTION 2
Which planets are most similar to Earth (Extended ESI approach)?
=====
```

METHODOLOGY:

1. Trained RF on non-ESI features → R² = 51% (ESI incomplete)
2. Trained RF on ESI + 10 non-ESI features → R² = 98.6%

3. Used Extended model predictions as Extended ESI
4. K-means clustering on 15 extended features
5. Ranked by distance to Earth in extended feature space

TOP 10 MOST SIMILAR PLANETS (Earth's cluster, Extended ESI):

Rank	Planet	Dist	Extended ESI	Original ESI	Orig Rank	Rank
Ext Rank	Rank Shift					
1	Earth	0.000	1.0000	0.9999	1	1
0						
2	Venus	1.597	0.8177	0.8415	19	2
+17						
3	Kepler-725 c	3.167	0.6705	0.5027	1145	3
+1142						
4	KOI-1831 d	0.899	0.6481	0.7669	76	4
+72						
5	Kepler-139 c	1.125	0.6351	0.4513	1649	5
+1644						
6	Kepler-452 b	7.048	0.6199	0.5513	738	6
+732						
7	Kepler-48 d	0.901	0.5982	0.5867	516	7
+509						
8	K2-18 b	7.253	0.5855	0.5731	592	8
+584						
9	Kepler-411 d	1.262	0.5832	0.4396	1785	9
+1776						
10	Kepler-260 c	1.171	0.5806	0.5395	830	10
+820						
=====						
=====						

```
[ ]: # H1: Popular planets hypothesis test
popular_planets_test = {
    "Kepler-452 b": "Earth's cousin",
    "TRAPPIST-1 e": "TRAPPIST HZ",
    "Proxima Cen b": "Closest to Solar System",
    "Kepler-186 f": "First Earth-size in HZ",
    "Kepler-442 b": "High ESI"
}

print("\n" + "="*100)
print(" "*30 + "HYPOTHESIS TEST (H1)")
print("Do popular 'Earth-like' planets rank highest in Extended ESI?")
print("=*100)
```

```

print(f"\n{'Planet':<20}{['Description':<30}{['Original Rank':<14}{['Proposed_Rank':<14}{['Shift':<10}]")
print(" "*100)

popular_ranks_proposed = []
for pname, desc in popular_planets_test.items():
    matches = df_imputed[df_imputed["pl_name"].str.contains(pname.split()[0], case=False, na=False)]
    if len(matches) > 0:
        p = matches.iloc[0]
        rank_orig = int(p['rank_esi'])
        rank_prop = int(p['rank_extended'])
        shift = rank_orig - rank_prop
        shift_str = f"+{shift}" if shift > 0 else str(shift)
        popular_ranks_proposed.append(rank_prop)
        print(f"{pname:<20}{desc:<30}{rank_orig:<14}{rank_prop:<14}{shift_str:<10}")
    else:
        print(f"{pname:<20}{desc:<30}{'NOT FOUND':<14}{'NOT FOUND':<14}{'-':<10}")

in_top_20 = sum(1 for r in popular_ranks_proposed if r <= 20)
print("\n" + " "*100)
print(f"RESULT: {in_top_20}/{len(popular_ranks_proposed)} popular planets in TOP 20 by Extended ESI")

if in_top_20 < len(popular_ranks_proposed) / 2:
    print("HYPOTHESIS SUPPORTED: Popular planets don't rank highest")
else:
    print("HYPOTHESIS NOT SUPPORTED: Popular planets do rank highly")
print("=="*100)

```

=====

=====

HYPOTHESIS TEST (H1)

Do popular 'Earth-like' planets rank highest in Extended ESI?

=====

=====

Planet	Description	Original Rank	Proposed Rank
Shift			

Kepler-452 b +732	Earth's cousin	738	6
----------------------	----------------	-----	---

TRAPPIST-1 e -117	TRAPPIST HZ	6	123
Proxima Cen b -1734	Closest to Solar System	265	1999
Kepler-186 f -1764	First Earth-size in HZ	235	1999
Kepler-442 b -1752	High ESI	247	1999

RESULT: 1/5 popular planets in TOP 20 by Extended ESI
HYPOTHESIS SUPPORTED: Popular planets don't rank highest

1.3.1 Final Visualizations

```
[29]: # Viz: Extended ESI vs Distance in Earth's cluster
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Left: Extended ESI vs Distance
axes[0].scatter(earth_cluster_ext_df['dist_earth_ext'], ↴
    earth_cluster_ext_df['extended_esi'],
    alpha=0.6, s=50, color='steelblue', label="Earth's cluster")
earth_data = earth_cluster_ext_df.iloc[0]
axes[0].scatter(earth_data['dist_earth_ext'], earth_data['extended_esi'],
    s=400, color='red', marker='*', edgecolors='black', ↴
    linewidths=2, label='Earth', zorder=10)
axes[0].set_xlabel('Distance to Earth (extended features)', fontsize=12)
axes[0].set_ylabel('Extended ESI', fontsize=12)
axes[0].set_title('Extended ESI vs Proximity in Earth\\'s Cluster', fontsize=14, ↴
    fontweight='bold')
axes[0].legend()
axes[0].grid(True, alpha=0.3)

# Right: Original vs Extended ESI
axes[1].scatter(df_imputed['esi'], df_imputed['extended_esi'], alpha=0.3, s=20, ↴
    color='gray')
axes[1].plot([0, 1], [0, 1], 'r--', linewidth=2, label='Perfect agreement')
axes[1].scatter(df_imputed.loc[earth_mask, 'esi'], df_imputed.loc[earth_mask, ↴
    'extended_esi'],
    s=400, color='red', marker='*', edgecolors='black', ↴
    linewidths=2, label='Earth', zorder=10)
axes[1].set_xlabel('Original ESI', fontsize=12)
axes[1].set_ylabel('Extended ESI', fontsize=12)
```

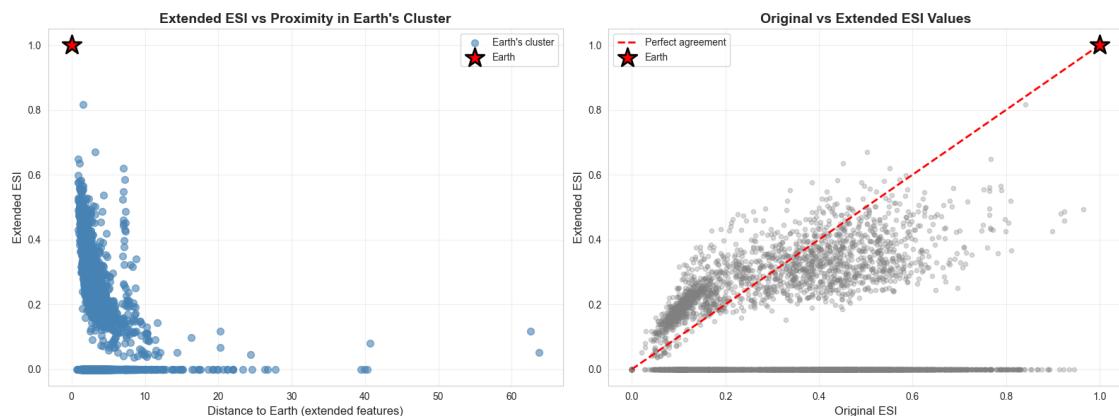
```

axes[1].set_title('Original vs Extended ESI Values', fontsize=14, u
    ↪fontweight='bold')
axes[1].legend()
axes[1].grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig(reports_dir / "extended_esi_final.png", dpi=150, u
    ↪bbox_inches="tight")
plt.show()

print("Saved: extended_esi_final.png")

```



Saved: extended_esi_final.png